



 Open in Colab

## ASSIGNMENT-3

**Problem1:**TitanicAssignmentSurvivalPrediction

```
In [1]: !pip install lime
```

```
Collecting lime
  Downloading lime-0.2.0.1.tar.gz (275 kB)
    ━━━━━━━━━━━━━━━━ 0.0/275.7 kB ? eta -:--:--
    ━━━━━━━━━━━━━━ 266.2/275.7 kB 13.6 MB/s eta 0:00:0
1
  ━━━━━━━━━━━━━━ 275.7/275.7 kB 6.9 MB/s eta 0:00:0
0
  Preparing metadata (setup.py) ... done
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (from lime) (3.10.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from lime) (2.0.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from lime) (1.16.1)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from lime) (4.67.1)
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.12/dist-packages (from lime) (1.6.1)
Requirement already satisfied: scikit-image>=0.12 in /usr/local/lib/python3.12/dist-packages (from lime) (0.25.2)
Requirement already satisfied: networkx>=3.0 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (3.5)
Requirement already satisfied: pillow>=10.1 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (11.3.0)
Requirement already satisfied: imageio!=2.35.0,>=2.33 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (2.37.0)
Requirement already satisfied: tifffile>=2022.8.12 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (2025.6.11)
Requirement already satisfied: packaging>=21 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (25.0)
Requirement already satisfied: lazy-loader>=0.4 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (0.4)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn>=0.18->lime) (1.5.1)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn>=0.18->lime) (3.6.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->lime) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib->lime) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib->lime) (4.59.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->lime) (1.4.9)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->lime) (3.2.3)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib->lime) (2.9.0.post0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7->matplotlib->lime) (1.17.0)
Building wheels for collected packages: lime
  Building wheel for lime (setup.py) ... done
  Created wheel for lime: filename=lime-0.2.0.1-py3-none-any.whl size=283834 sha256=fda785971bd3c4308614a80a827f3b7395a673cf112456a40f575b86c017b2a8
```

```
Stored in directory: /root/.cache/pip/wheels/e7/5d/0e/4b4fff9a47468fed5633211  
fb3b76d1db43fe806a17fb7486a  
Successfully built lime  
Installing collected packages: lime  
Successfully installed lime-0.2.0.1
```

```
In [2]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import LabelEncoder  
from sklearn.ensemble import GradientBoostingClassifier  
from sklearn.metrics import classification_report, accuracy_score, ConfusionMatrix
```

```
In [3]: import lime  
import lime.lime_tabular
```

```
In [4]: df = pd.read_csv("/content/titanic dataset.csv")  
print("Dataset shape:", df.shape)  
df.head()
```

Dataset shape: (891, 12)

```
Out[4]:   PassengerId  Survived  Pclass      Name     Sex   Age  SibSp  Parch  Ticket  
          0            1        0       3  Braund,  
                         Mr. Owen  
                         Harris  male  22.0      1      0    A/5  
                         21171  
          1            2        1       1  Cumings,  
                         Mrs. John  
                         Bradley  
                         (Florence  
                         Briggs  
                         Th...  female  38.0      1      0    PC  
                         17599  
          2            3        1       3  Heikkinen,  
                         Miss.  
                         Laina  female  26.0      0      0  STON/  
                         O2  
                         3101282  
          3            4        1       1  Futrelle,  
                         Mrs.  
                         Jacques  
                         Heath  
                         (Lily May  
                         Peel)  female  35.0      1      0  113803  
          4            5        0       3  Allen, Mr.  
                         William  
                         Henry  male  35.0      0      0  37345C
```

```
In [5]: df['Age'] = df['Age'].fillna(df['Age'].median())  
df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0])
```

```
In [6]: label_cols = ['Sex', 'Embarked']
```

```
le = LabelEncoder()
for col in label_cols:
    df[col] = le.fit_transform(df[col])
```

```
In [7]: features = ['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked']
X = df[features]
y = df['Survived']
```

```
In [8]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [9]: model = GradientBoostingClassifier(random_state=42)
model.fit(X_train, y_train)
```

```
Out[9]: GradientBoostingClassifier
```

```
GradientBoostingClassifier(random_state=42)
```


```
In [10]: y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.8100558659217877
          precision    recall  f1-score   support

           0       0.81      0.89      0.85      105
           1       0.81      0.70      0.75       74

    accuracy                           0.81      179
   macro avg       0.81      0.79      0.80      179
weighted avg       0.81      0.81      0.81      179
```

```
In [11]: ConfusionMatrixDisplay.from_predictions(y_test, y_pred, cmap="Blues")
plt.show()
```




```
In [12]: explainer = lime.lime_tabular.LimeTabularExplainer(  
    training_data = np.array(X_train),  
    feature_names = features,  
    class_names = ['Died', 'Survived'],  
    mode='classification'  
)
```

```
In [13]: i = 10  
exp = explainer.explain_instance(X_test.iloc[i], model.predict_proba, num_feat  
exp.show_in_notebook(show_table=True)
```

```
/usr/local/lib/python3.12/dist-packages/lime/discretize.py:110: FutureWarning:  
Series.__getitem__ treating keys as positions is deprecated. In a future version,  
integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`  
    ret[feature] = int(self.lambdas[feature](ret[feature]))  
/usr/local/lib/python3.12/dist-packages/lime/discretize.py:110: FutureWarning:  
Series.__setitem__ treating keys as positions is deprecated. In a future version,  
integer keys will always be treated as labels (consistent with DataFrame behavior). To set a value by position, use `ser.iloc[pos] = value`  
    ret[feature] = int(self.lambdas[feature](ret[feature]))  
/usr/local/lib/python3.12/dist-packages/lime/lime_tabular.py:544: FutureWarning:  
Series.__getitem__ treating keys as positions is deprecated. In a future version,  
integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`  
    binary_column = (inverse_column == first_row[column]).astype(int)  
/usr/local/lib/python3.12/dist-packages/sklearn/utils/validation.py:2739: UserWarning:  
X does not have valid feature names, but GradientBoostingClassifier was  
fitted with feature names  
    warnings.warn(  
/usr/local/lib/python3.12/dist-packages/lime/discretize.py:110: FutureWarning:  
Series.__getitem__ treating keys as positions is deprecated. In a future version,  
integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`  
    ret[feature] = int(self.lambdas[feature](ret[feature]))  
/usr/local/lib/python3.12/dist-packages/lime/discretize.py:110: FutureWarning:  
Series.__setitem__ treating keys as positions is deprecated. In a future version,  
integer keys will always be treated as labels (consistent with DataFrame behavior). To set a value by position, use `ser.iloc[pos] = value`  
    ret[feature] = int(self.lambdas[feature](ret[feature]))  
/usr/local/lib/python3.12/dist-packages/lime/lime_tabular.py:427: FutureWarning:  
Series.__getitem__ treating keys as positions is deprecated. In a future version,  
integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`  
    discretized_instance[f])]
```

```
In [14]: proba = model.predict_proba(X_test)[:,1]  
plt.hist(proba, bins=20, color='skyblue', edgecolor='black')  
plt.xlabel("Predicted probability of survival")  
plt.ylabel("Number of passengers")  
plt.title("Survival Probability Distribution")  
plt.show()
```



## ANALYSIS REPORT:

The Titanic dataset is used to predict whether a passenger survived or not. For this task, we used a machine learning model called Gradient Boosting. The dataset contains details such as passenger age, gender, class, number of family members, and fare. Before training the model, we cleaned the data by filling in missing values and converting text features (like male/female and boarding port) into numbers.

Once trained, the model gave good accuracy in predicting survival. The important factors that influenced survival were gender, class, and age. Women, children, and people in higher classes had a better chance of survival compared to men and passengers in lower classes.

To understand individual predictions, we applied LIME (Local Interpretable Model-Agnostic Explanations). LIME explains why the model made a certain prediction by highlighting which features contributed positively or negatively. For example, in one case, being female and traveling in a higher class increased the survival probability, while being male and older reduced it.

We also plotted the predicted probabilities of survival. This showed a clear

difference: passengers with high chances of survival were mostly women and upper-class travelers, while those with low chances were mostly men in lower classes.

## Problem2:DiabetesPrediction

```
In [15]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, ConfusionMatrix

import lime
import lime.lime_tabular
```

```
In [16]: df = pd.read_excel("/content/diabetes.csv.xlsx")
print("Shape:", df.shape)
df.head()
```

Shape: (768, 9)

```
Out[16]:    Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  DiabetesF
0            6        148             72           35         0   33.6
1            1         85             66           29         0   26.6
2            8        183             64           0         0   23.3
3            1         89             66           23        94   28.1
4            0        137             40           35        168  43.1
```

```
In [17]: X = df.drop('Outcome', axis=1)    # Outcome: 1 = Diabetic, 0 = Non-diabetic
y = df['Outcome']
```

```
In [18]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [19]: X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)
```

```
In [20]: model = LogisticRegression(max_iter=500)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
```


```
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.7532467532467533
          precision    recall  f1-score   support

           0       0.81      0.80      0.81      99
           1       0.65      0.67      0.66      55

    accuracy                           0.75      154
   macro avg       0.73      0.74      0.73      154
weighted avg       0.76      0.75      0.75      154
```


```
In [21]: ConfusionMatrixDisplay.from_predictions(y_test, y_pred, cmap="Blues")
plt.show()
```



```
In [22]: explainer = lime.lime_tabular.LimeTabularExplainer(
    training_data=X_train,
    feature_names=X.columns,
    class_names=['Non-Diabetic', 'Diabetic'],
    mode='classification'
)
```

```
In [23]: i = 5
exp = explainer.explain_instance(X_test[i], model.predict_proba, num_features=
exp.show_in_notebook(show_table=True)
```

```
In [24]: proba = model.predict_proba(X_test)[:,1]
plt.hist(proba, bins=20, color='lightcoral', edgecolor='black')
plt.xlabel("Predicted probability of Diabetes")
plt.ylabel("Number of Patients")
plt.title("Probability Distribution for Diabetes Risk")
plt.show()
```



#### INTERPRET RESULT

The Logistic Regression model predicts whether a person is diabetic or not based on health features like glucose, BMI, age, and pregnancies. The model gave good accuracy. LIME showed that high glucose and high BMI increase diabetes risk, while low values push towards non-diabetic. This helps doctors understand both the prediction and the main medical risk factors